

# IAPR 2024 - Project

## Coin Detection Challenge

Theo Heng, Océane Voland & Maïka Nogarotto



# Project overview

## General structure

- **Segmentation**
  - Background detection
  - Coins recognition
  - Coins isolation
- **Feature analysis**
  - Area extraction
  - Contours enhancement
- **Classification**
  - **Template matching**
  - **Deep learning model**



Illustration created by Olivier Lemer, a former student of the IC faculty. (olivier.lemer@alumni.epfl.ch)

Coin 1



Coin 2



Coin 3



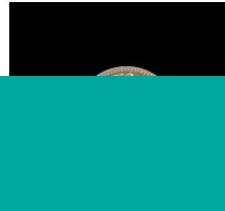
Coin 6



Coin 7



Coin 8



Coin 11



Coin 12

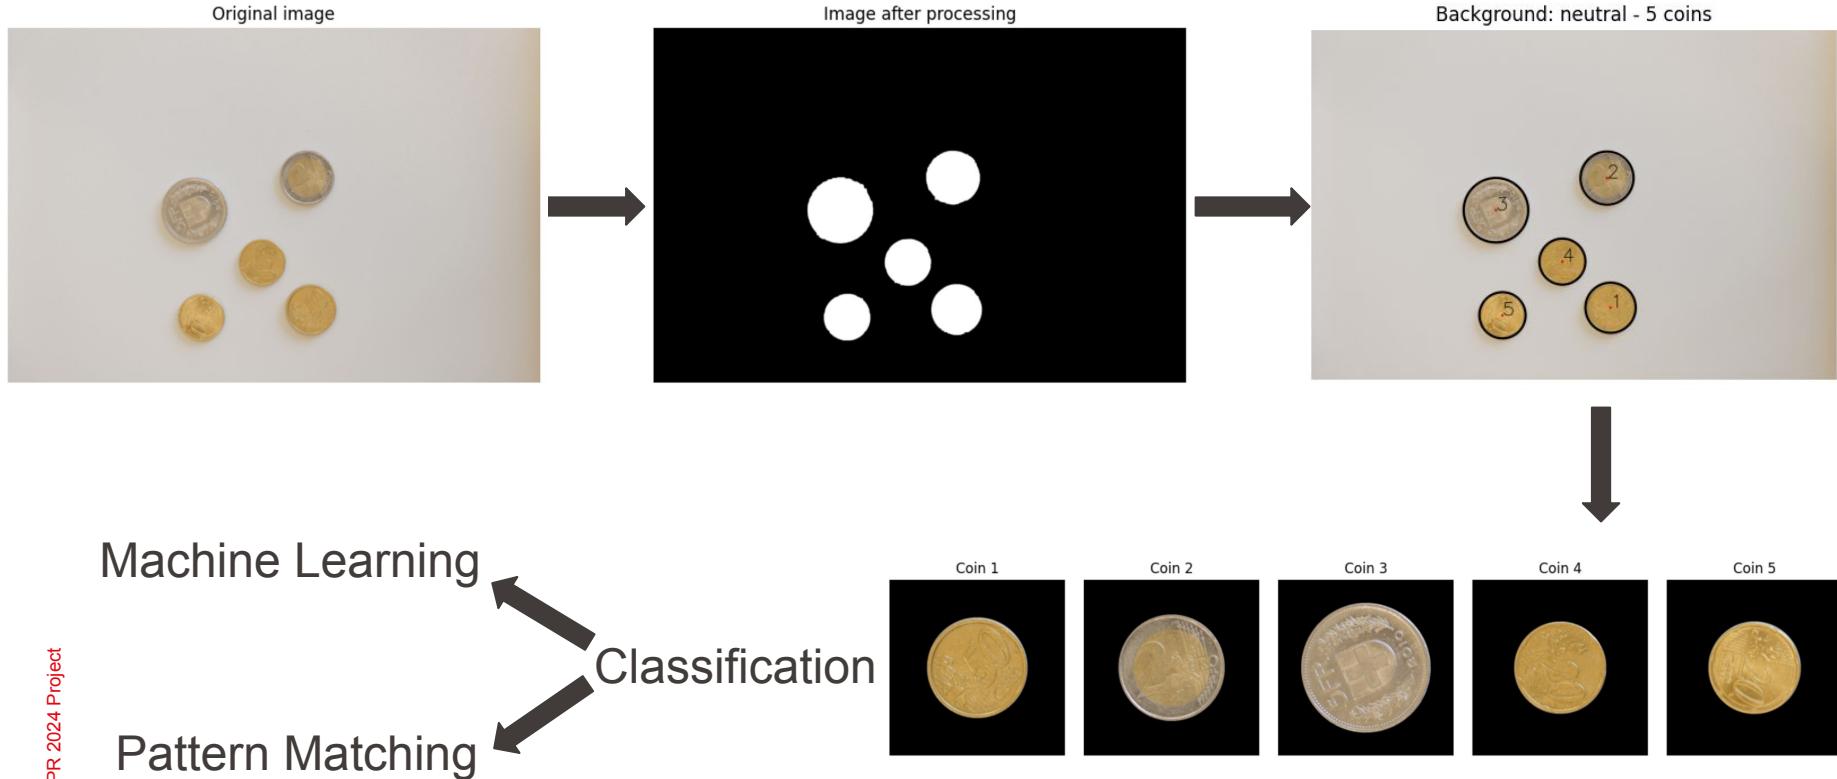


# 1. Segmentation



# 1. Segmentation

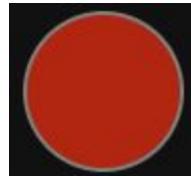
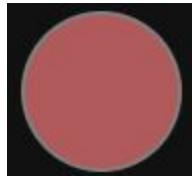
## General pipeline



# 1. Segmentation

## Background detection

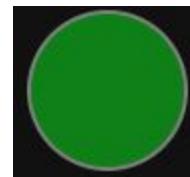
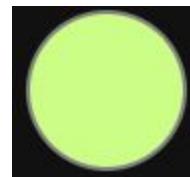
- Type of background detected using thresholds on specific colors



A lot of pink/red



Hand background



A lot of blue/green



Noisy background



Else



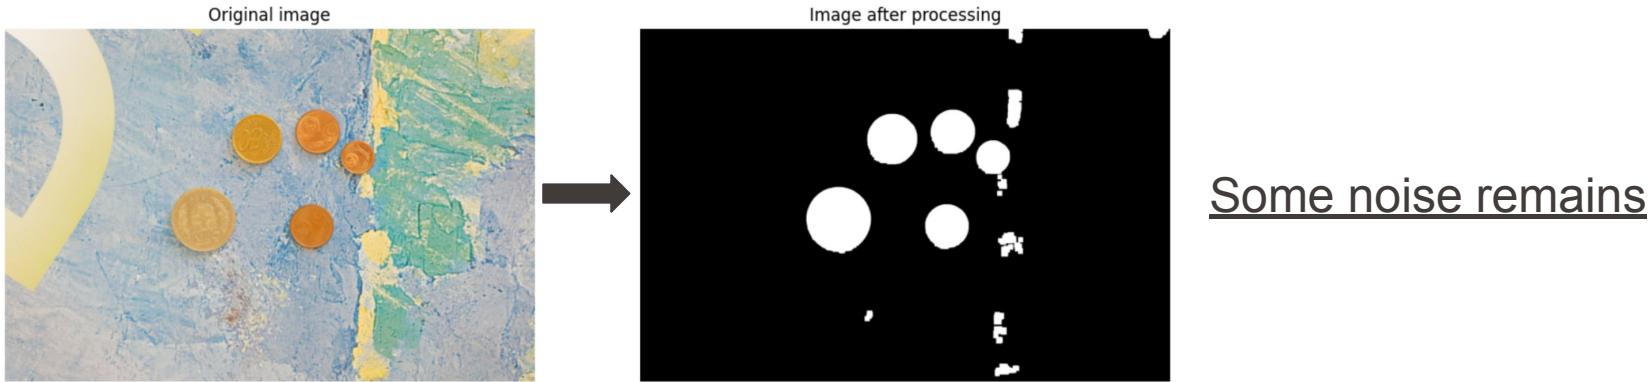
Neutral background



# 1. Segmentation

## Segmentation

- **Blurring** (Median, Gaussian)
- **Thresholding** to binarize the image into foreground and background
  - Thresholds depend on background and are based on coin's colors
- **Morphological operations**
  - Closing : remove holes
  - Opening : get smoother contours

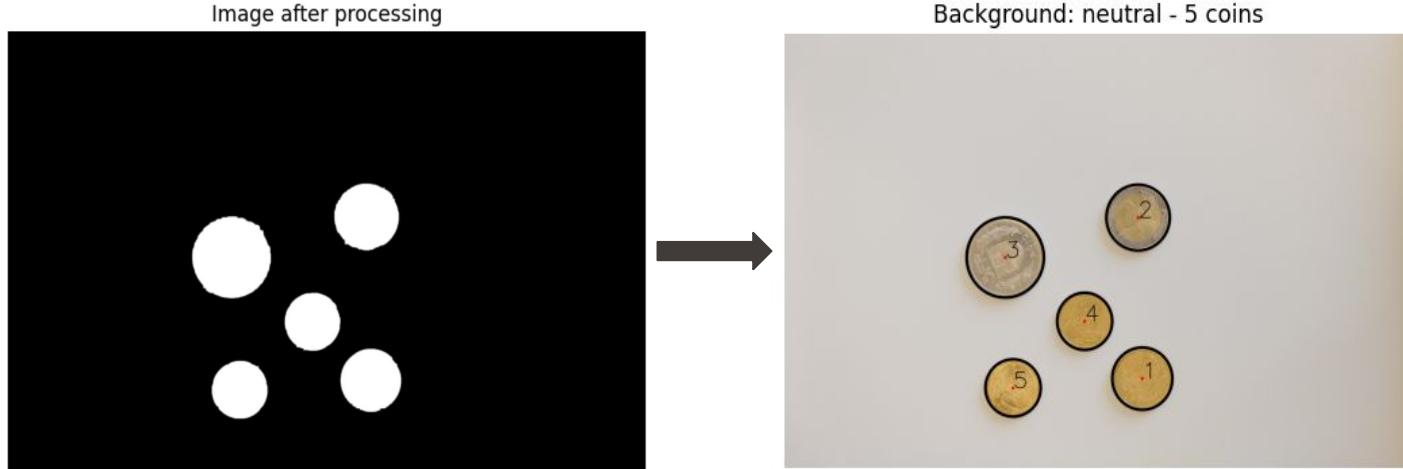


# 1. Segmentation

## Coin identification

- **Find circles in images ( cv2.HoughCircles )**

- Optimal parameters ensure that noise is not considered as coins
- Parameters depend on background
- Returns center and radius of each circle

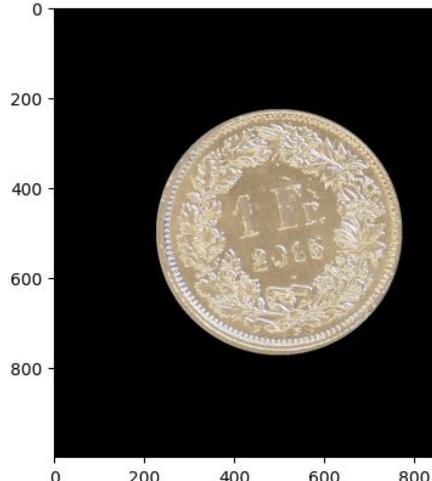
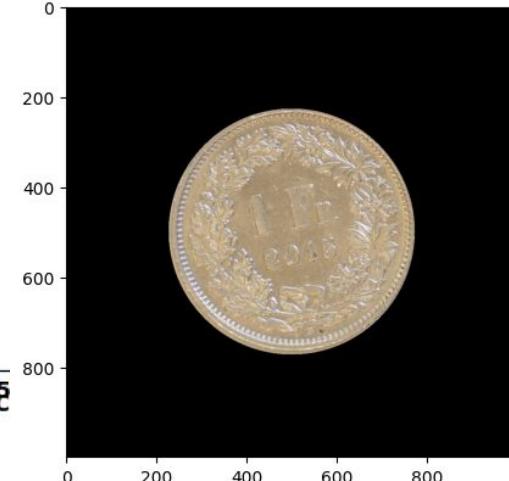
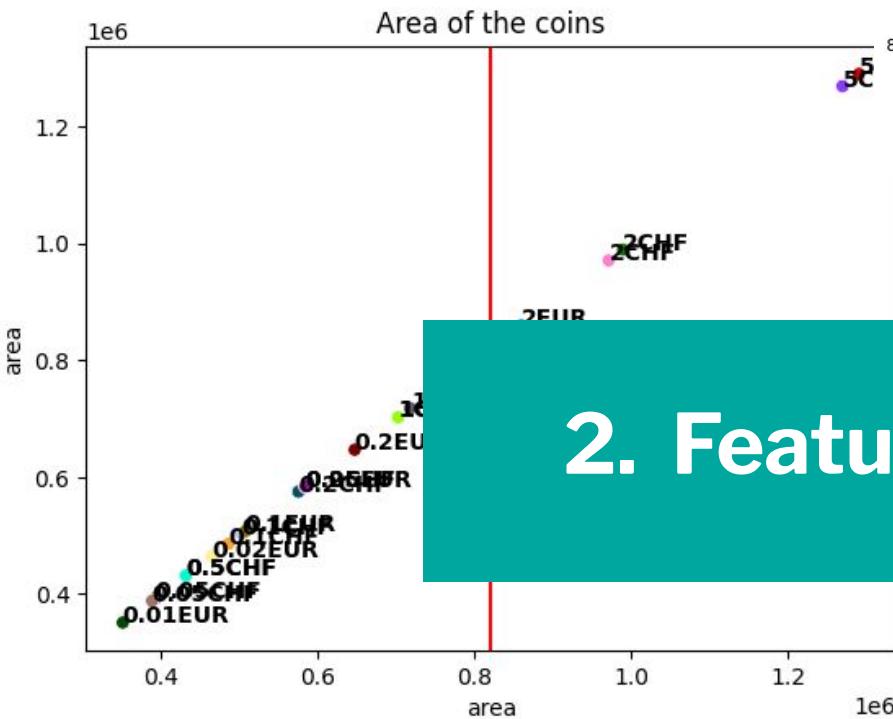


# Segmentation

## Coin isolation

- Using the detected circles, we can create sub-images containing individual coin:
  - 1000x1000 pixels
  - Black background
  - Coin centered
- Segmentation process **100% automatic**
- **100% of the coins detected**
- **Not robust** to external images to the dataset



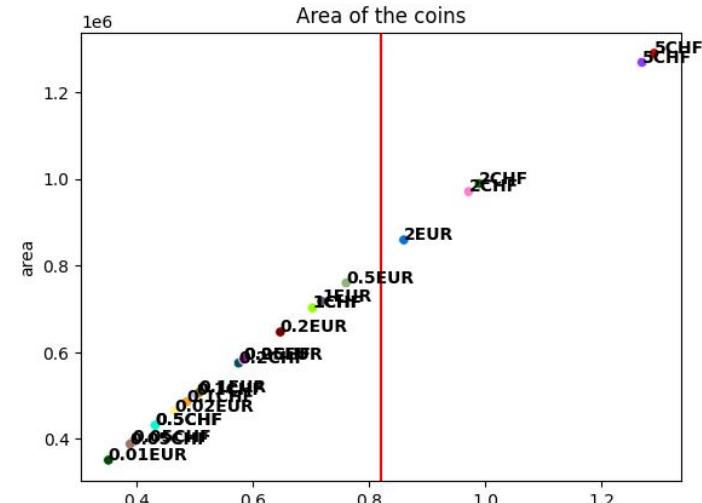


## 2. Features Analysis

# 2. Features Analysis

## Area extraction

- **Get the number of non-black pixels on segmented images**
  - Faster way than computing the radius again
- **Can analysis the classes depending on their area**
  - Can be used to distinguish small from big coins
    - Useful to save computation resources and time
  - Not meaningful enough to be used alone for classification
    - Small coins have too similar area



# 2. Features Analysis

## Contour extraction and enhancement

- Contours extracted using Canny edge detection

- Use histogram equalisation to enhance the contrast
- Find the contours

- Contour enhancement

- Get the images with equalized detected edges
- Increase the value of the pixel in the image at the contours location

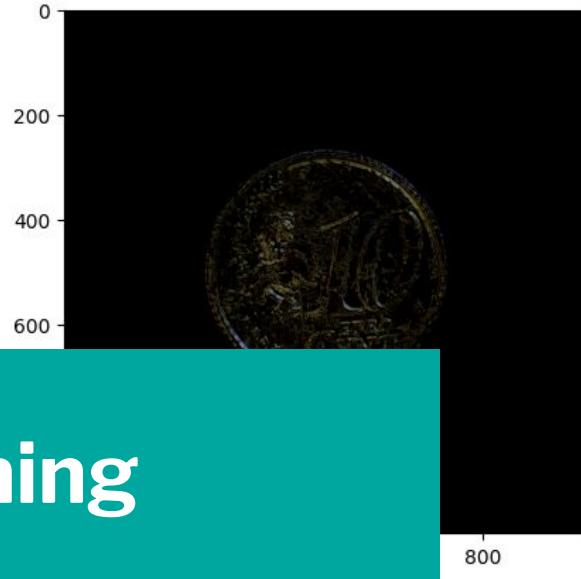
- Benefits

- Allows a better differentiation of the coins





## 3. Template Matching



# 3. Template Matching

## Principle and concepts

- **Method in general**

- Set an image as a template
- Find on the image where the template matches the best

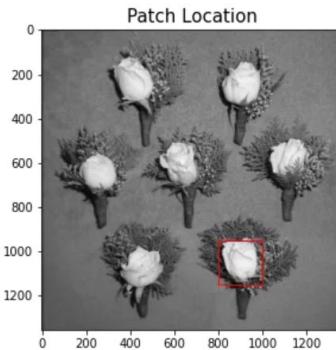


Figure 2: Template/Patch Image (Image by Author)

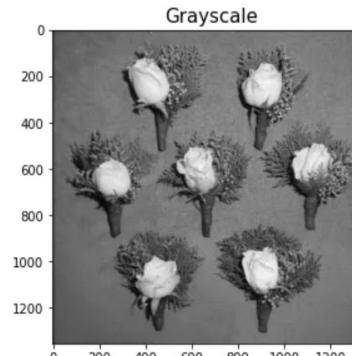
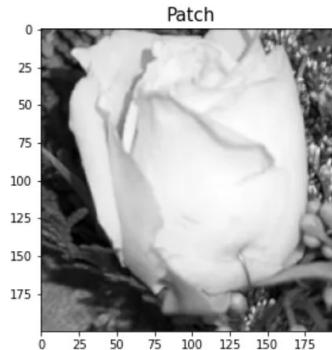
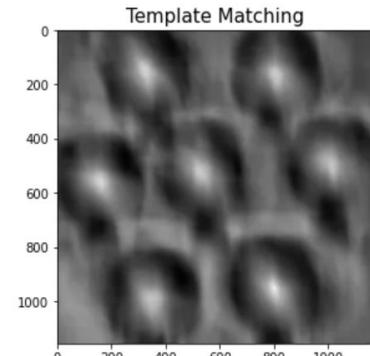


Figure 3: Performing Template Matching (Image by Author)



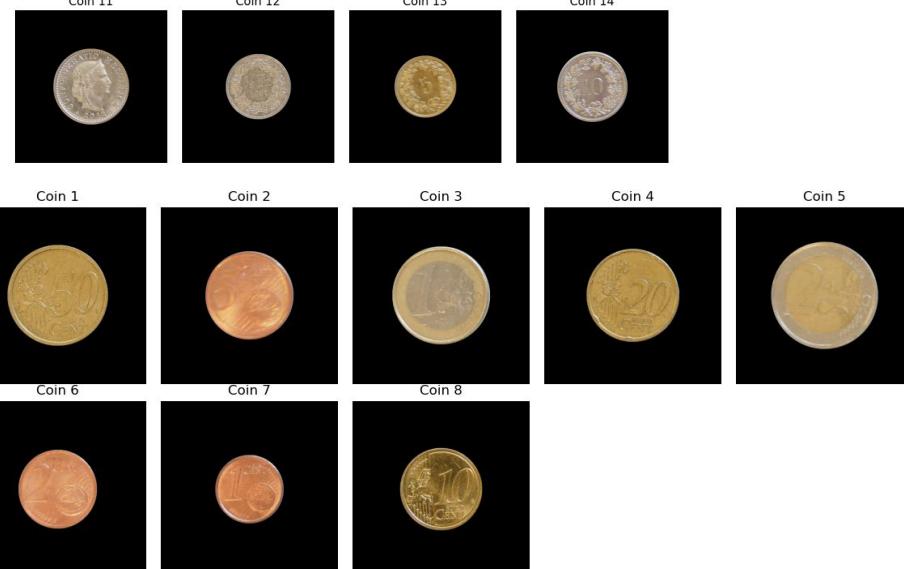
# 3. Template Matching

## Our adaptation

- Which reference coin is our testing coin the closest ?
  - Compare our testing coin to each reference coin and take the most similar one



Reference coins



# 3. Template Matching

## Application

- Methods used

- Squared difference



Test

VS



Reference

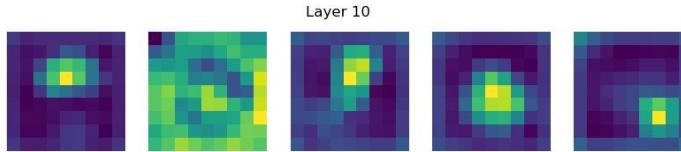


# 3. Template Matching

## Results

- **Accuracy : 66,24%**
- **Advantages**
  - Only need one example of each coins type
  - Don't need training
  - Fast
  - Transparent pipeline
- **Drawbacks**
  - Lack of precision
  - Outperformed by most ML technics
  - Struggle with outliner coins
  - Cropping imprecisions can dramatically impact the result





Layer 10



Layer 130

## 4. Machine Learning

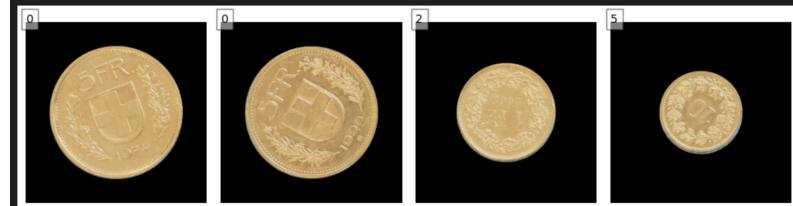
google

/efficientnet-b3



huggingface.co

Found 4 coins  
Total: 0 €, 11.1 CHF, and 0 000



# 4. Machine Learning

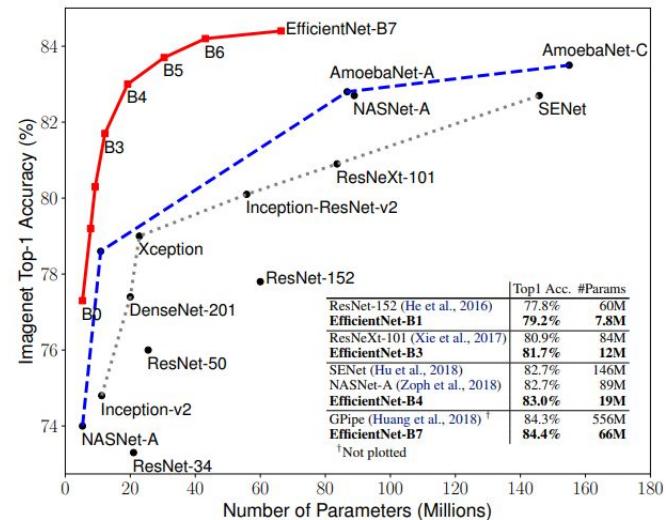
## Motivations

- **Why a pretrained model?**

- Reduced Training Time and Computational Resources
- Better Generalization and Robust Feature Extraction
- Effective Use of Limited Data

- **Why EfficientNet B3?**

- Lightweight and SOTA performances
- Well documented and easy to use
- Nice trade-off performance/size



# 4. Machine Learning

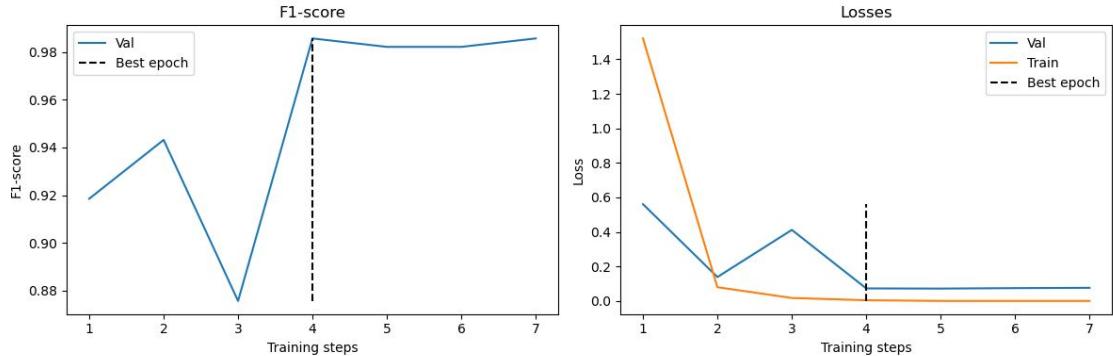
## Training specifications

- **Dataset**
  - Data Augmentation -> 6000 train images - 800 validation images
- **Special features**
  - Adaptive learning rate with scheduler
  - Gradient accumulation to emulate bigger batch sizes without RAM overflow
  - Confusion Matrix to identify classification flaws

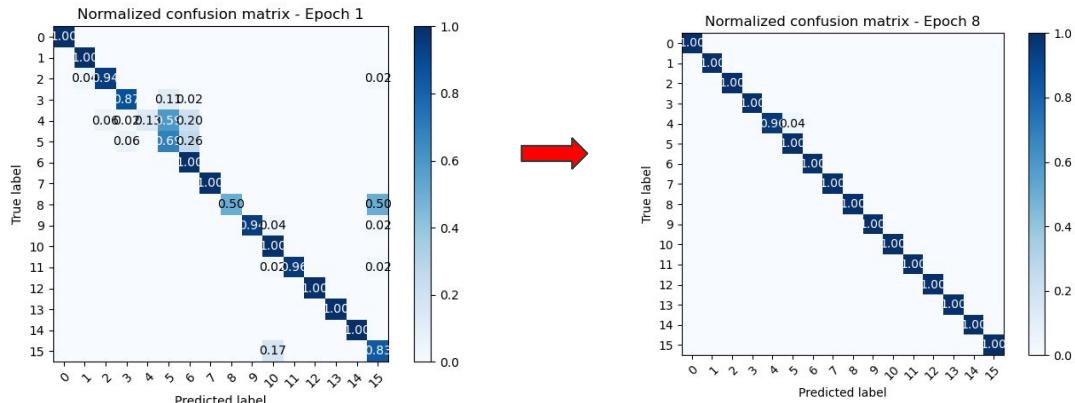
|                 | Epochs | Batch size with accumulation | Learning Rate | Training time |
|-----------------|--------|------------------------------|---------------|---------------|
| EfficientNet B3 | 10     | $32 \times 2 = 64$           | 5e-5          | 60 mins       |

# 4. Machine Learning Results

- Training loss:



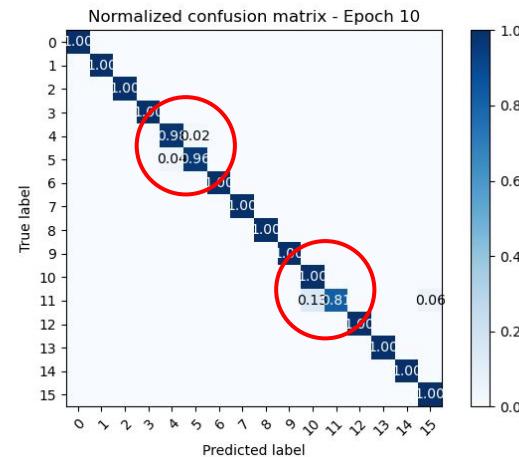
- Confusion matrices:



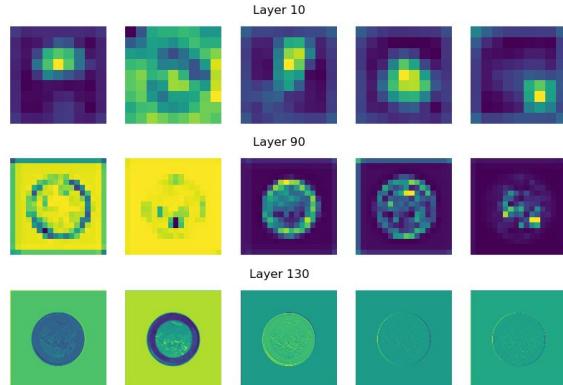
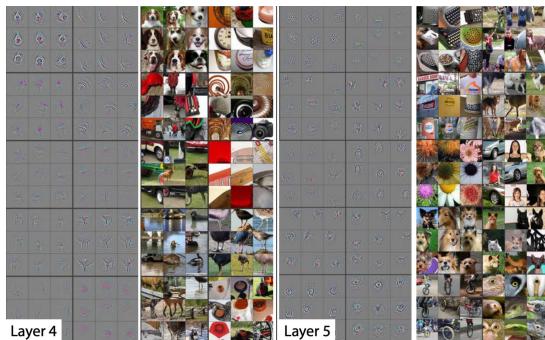
- Results: **98.4%**

# 4. Machine Learning Analysis

- Why is 100% so hard to reach ?
  - 20 and 10 cents often confused together
  - Not enough data to fix this



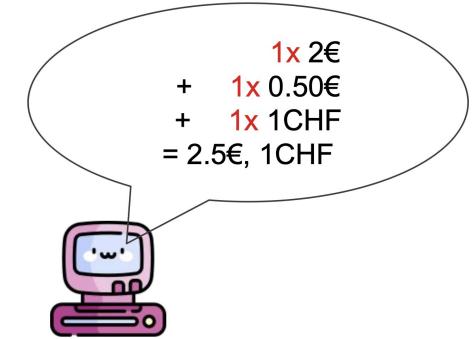
- What are the features extracted by the model?



Example from EE-451 slides



## 5. Conclusion



# 5. Conclusion

## Project summary

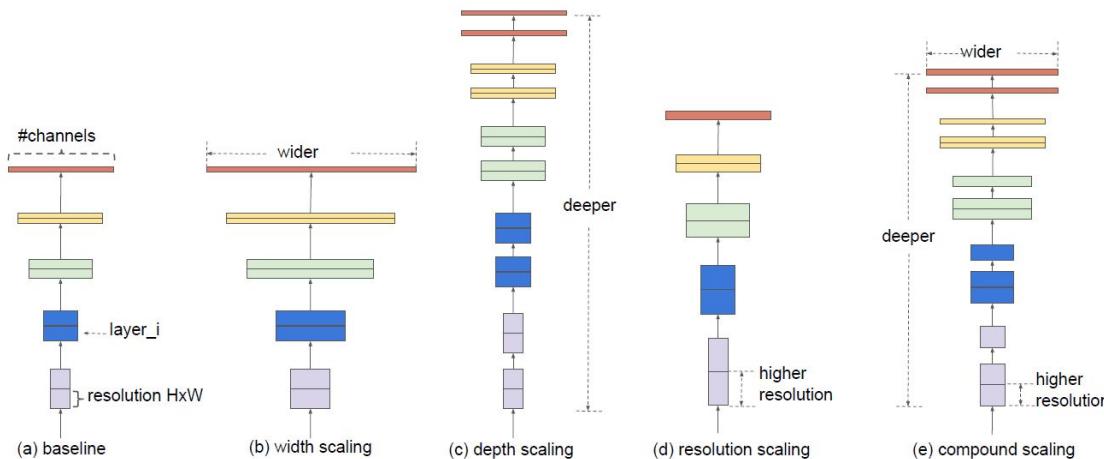
- **Segmentation:** 100% of coins detected and extracted
- **Template matching:** 66.24%
  - Require very few data and no training
- **Deep learning model:** 98.4%
  - Require training and manual annotation



# 6. Annexe

## EfficientNet B0 -> EfficientNet B3

|                 | NB parameters | Resolution | Model size | Output channels |
|-----------------|---------------|------------|------------|-----------------|
| EfficientNet B0 | 5.3 M         | 224 x 224  | 16.4 M     | 1280            |
| EfficientNet B3 | 12 M          | 300 x 300  | 43.4 MB    | 1536            |

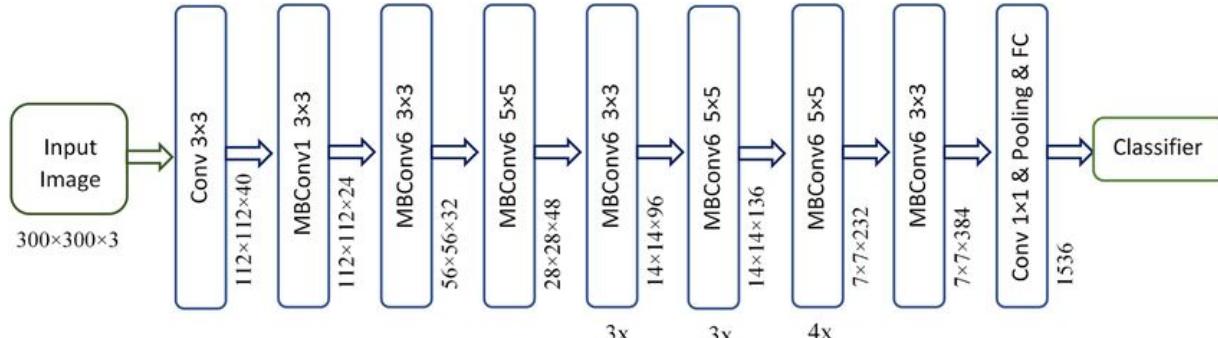


**Compound Scaling:**  
scaling factor to depth, width, and resolution based on a compound coefficient ( $\varphi$ ).  
-> ensures that the network grows in a balanced manner.

# 6. Annexe

## EfficientNet backbone

$$\text{swish}(x) = x \text{ sigmoid}(\beta x) = \frac{x}{1 + e^{-\beta x}}.$$



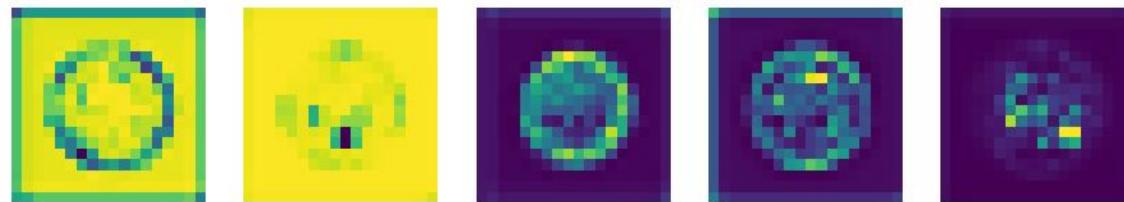
- **1st convolution layer of 32 3x3 filters -> 32 features map**
- **Successive sets of MBConv blocks**
  - Expansion, Depthwise convolution, projection, skip connection, batch normalisation+activation
- Each set increases the number of filters and may change the spatial dimensions by downsampling
- **Global average pooling layer ->** average each feature map to a single dimension
- **Fully connected (dense) layer ->** combines the features extracted to produce final logits
- **Output layer ->** Use softmax to produce the final class predictions

# 6. Annexe

## Extracting features from the model

- **Hook Function:** Captures outputs of convolutional layers.
- Iterates over all layers in the model and registers the hook function to each convolutional layer
- Disables gradient calculation to save memory and computations, as we are only interested in the outputs, not in training.
- **Forward Pass:** Passes the input image through the model to gather feature maps.
- Removes all the hooks after the forward pass to avoid any side effects on the model.

Layer 90



# 6. Annexe

## Choice of Optimizer

- **ADAM (Adaptive Moment Estimation)**
  - uses estimates of first and second moments of the gradients (mean and uncentered variance) to adapt the learning rate for each parameter
  - computationally efficient and requires little memory
- **ADAMW**
  - improvement over Adam by decoupling weight decay from the gradient update. Theoretical better convergence
  - Requires careful tuning of the weight decay parameter. The improvement in performance might not be significant for all tasks compared to standard Adam
- **SVG (Stochastic average gradient)**
  - Reduces variance of gradient estimates by averaging gradients from all past iterations. Can converge faster than SGD in some cases
  - Requires storing the entire gradient history -> higher memory usage. Not as widely used or well-understood as Adam.

# 2. Features analysis

## Overview

- **Area extraction**

- Get the number of non-black pixels
- Allows to separate big from small coins

- **Contours enhancement**

- Get the edge using Canny edge detection
- Deep learning model

# 3. Template matching

## Overview

- **Principle and concepts**
  - Method in general
  - Our adaptation
- **Application**
  - Require data
  - Methods used
- **Results**
  - Advantages
  - Drawbacks
- **Possible improvements**

# 3. Template matching

## Application

- **Data required**
  - Only requires the reference images



# 3. Template matching

## Possible improvements

- **Create an outlier “database”**
  - Require to use the training images
- **Add more images to reference sets**
  - Require more computational resources and time
- **Rotate the image during the comparison**
  - Require more computational resources and time
- **Get more general reference images**
  - Fuse all coin of the same time together
  - Use SVD or KLT to get eigenvectors
  - Require to use the training images
- **Use other comparison metrics**